

Report 4
Introduction to CUDA and OpenCL
SM, memory prefetching and nvprof.

Dariusz Biela

Karol Sawicki

1. Introduction

At the time of our fourth lab classes we were trying to gain more performance with using Unified Memory. To do this we had to learn about memory prefetching mechanism.

Memory prefetching is asynchronous mechanism that sends data from the device to the CPU before it needs the data. This allows to do operations quicker, because when CPU needs the data it already is available. Memory prefetching works in the opposite direction too.

Memory prefetching basically allows hardware to avoid page fault problems. Page fault is a type of hardware exception, that can be occurred when program tries to access memory that is not currently mapped.

As an addition we learned about SM – Streaming Multiprocessors – every nvidia GPU chip is built of them.

2. Calculations and measurements

To start investigations we tried to use nvprof tool with the sample codes. We decided to show and discuss our results:

```
==31373== NVPROF is profiling process 31373, command: ./page_faultsCPU.exe
==31373== Profiling application: ./page_faultsCPU.exe
==31373== Profiling result:
No kernels were profiled.

```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
API calls:	97.41%	251.89ms	1	251.89ms	251.89ms	251.89ms	cudaMallocManaged
	1.92%	4.9585ms	1	4.9585ms	4.9585ms	4.9585ms	cudaFree
	0.36%	929.66us	194	4.7920us	907ns	176.14us	cuDeviceGetAttribute
	0.27%	688.50us	2	344.25us	240.53us	447.96us	cuDeviceTotalMem
	0.03%	85.206us	2	42.603us	39.879us	45.327us	cuDeviceGetName
	0.01%	22.488us	2	11.244us	6.9140us	15.574us	cuDeviceGetPCIBusId
	0.00%	5.7970us	3	1.9320us	978ns	2.5840us	cuDeviceGetCount
	0.00%	5.5180us	4	1.3790us	1.1170us	1.9560us	cuDeviceGet
	0.00%	2.5840us	2	1.2920us	1.1870us	1.3970us	cuDeviceGetUuid

```
==31373== Unified Memory profiling result:
Total CPU Page faults: 384
```

Picture 1 CPU only page faults

```

==31394== NVPROF is profiling process 31394, command: ./page_faultsCPUGPU.exe
==31394== Profiling application: ./page_faultsCPUGPU.exe
==31394== Profiling result:
   Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 100.00%    65.840ms    1    65.840ms    65.840ms    65.840ms    deviceKernel(int*, int)
  API calls:  76.88%    249.33ms    1    249.33ms    249.33ms    249.33ms    cudaMallocManaged
              20.30%    65.839ms    1    65.839ms    65.839ms    65.839ms    cudaDeviceSynchronize
              2.23%    7.2402ms    1    7.2402ms    7.2402ms    7.2402ms    cudaFree
              0.28%    906.54us    194    4.6720us    907ns    163.99us    cuDeviceGetAttribute
              0.21%    689.54us    2    344.77us    240.74us    448.80us    cuDeviceTotalMem
              0.05%    164.69us    1    164.69us    164.69us    164.69us    cudaLaunchKernel
              0.03%    88.559us    2    44.279us    40.508us    48.051us    cuDeviceGetName
              0.01%    22.070us    2    11.035us    6.2850us    15.785us    cuDeviceGetPCIBusId
              0.00%    5.5170us    4    1.3790us    978ns    2.0250us    cuDeviceGet
              0.00%    5.0980us    3    1.6990us    977ns    2.3750us    cuDeviceGetCount
              0.00%    2.5840us    2    1.2920us    1.1170us    1.4670us    cuDeviceGetUuid

==31394== Unified Memory profiling result:
Device "GeForce RTX 2060 (0)"
   Count  Avg Size  Min Size  Max Size  Total Size  Total Time  Name
    5176  12.492KB  4.0000KB  128.00KB  63.14844MB  12.88730ms  Host To Device
Total CPU Page faults: 384

```

Picture 2 CPU to GPU page faults

```

==31342== NVPROF is profiling process 31342, command: ./page_faultsGPU.exe
==31342== Profiling application: ./page_faultsGPU.exe
==31342== Profiling result:
   Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 100.00%    35.703ms    1    35.703ms    35.703ms    35.703ms    deviceKernel(int*, int)
  API calls:  87.01%    264.75ms    1    264.75ms    264.75ms    264.75ms    cudaMallocManaged
              11.73%    35.678ms    1    35.678ms    35.678ms    35.678ms    cudaDeviceSynchronize
              0.61%    1.8535ms    1    1.8535ms    1.8535ms    1.8535ms    cudaFree
              0.31%    943.63us    194    4.8640us    907ns    170.41us    cuDeviceGetAttribute
              0.24%    729.98us    2    364.99us    254.15us    475.83us    cuDeviceTotalMem
              0.06%    189.90us    1    189.90us    189.90us    189.90us    cudaLaunchKernel
              0.03%    91.491us    2    45.745us    41.276us    50.215us    cuDeviceGetName
              0.01%    23.886us    2    11.943us    5.7270us    18.159us    cuDeviceGetPCIBusId
              0.00%    5.4480us    4    1.3620us    978ns    2.1650us    cuDeviceGet
              0.00%    5.2390us    3    1.7460us    1.1880us    2.3050us    cuDeviceGetCount
              0.00%    2.3740us    2    1.1870us    1.1870us    1.1870us    cuDeviceGetUuid

==31342== Unified Memory profiling result:
Device "GeForce RTX 2060 (0)"
   Count  Avg Size  Min Size  Max Size  Total Size  Total Time  Name
    591    -    -    -    -    35.27318ms  Gpu page fault groups

```

Picture 3 GPU only page faults

```

[cuda-lab07@lhcbgpu2 lab-02042020]$ nvprof ./page_faultsGPUCPU.exe
==31415== NVPROF is profiling process 31415, command: ./page_faultsGPUCPU.exe
==31415== Profiling application: ./page_faultsGPUCPU.exe
==31415== Profiling result:
   Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 100.00%    30.370ms    1    30.370ms    30.370ms    30.370ms    deviceKernel(int*, int)
  API calls:  88.19%    310.08ms    1    310.08ms    310.08ms    310.08ms    cudaMallocManaged
              8.64%    30.367ms    1    30.367ms    30.367ms    30.367ms    cudaDeviceSynchronize
              2.63%    9.2428ms    1    9.2428ms    9.2428ms    9.2428ms    cudaFree
              0.26%    918.97us    194    4.7360us    907ns    179.28us    cuDeviceGetAttribute
              0.20%    687.59us    2    343.79us    240.32us    447.26us    cuDeviceTotalMem
              0.06%    194.65us    1    194.65us    194.65us    194.65us    cudaLaunchKernel
              0.03%    89.745us    2    44.872us    39.320us    50.425us    cuDeviceGetName
              0.01%    21.720us    2    10.860us    6.0060us    15.714us    cuDeviceGetPCIBusId
              0.00%    5.2390us    4    1.3090us    908ns    2.0260us    cuDeviceGet
              0.00%    5.2380us    3    1.7460us    978ns    2.5140us    cuDeviceGetCount
              0.00%    2.4440us    2    1.2220us    1.1170us    1.3270us    cuDeviceGetUuid

==31415== Unified Memory profiling result:
Device "GeForce RTX 2060 (0)"
   Count  Avg Size  Min Size  Max Size  Total Size  Total Time  Name
    768  170.67KB  4.0000KB  0.9961MB  128.0000MB  11.08192ms  Device To Host
    585    -    -    -    -    29.99165ms  Gpu page fault groups
Total CPU Page faults: 384

```

Picture 4 GPU to CPU page faults

Pictures from 1 to 4 are describing nvprof results when exemplary code was profiled. We can see that CPU page faults are always the same – when they are. The more interesting are GPU page faults – in each case number of page fault groups is different and every time we see time that we are losing with these page faults. This lost time could be saved by good memory prefetching use.

3. Conclusions

Our nvprof tool shows us that page faults problems are real world problems and – when they are – we could lose so much time. The *Picture 3* shows that we lost 35.3 ms – that may seem really small amount of lost time, but when we look at the kernel execution time - 35.7 ms – we know that it is really problematic, especially when we have to execute kernel hundreds of thousands times.

Usage of memory prefetching can be helpful to avoid page faults, especially when we know when data is needed by the host or nvidia device that we are using. In the past, with older GPU chips every operation was ended with memory synchronization even if that was not necessary. Today we can use prefetching with unified memory only when we really have to synchronize our data and avoid page faults, that cause time loss.